




Original word-task run, fluke candidate (just for demonstration) | compositional:gerund_first

allenai/Olmo-3-1025-7B | wrong no explicit context | Corruption: not recorded | Mode: random | Scope: not recorded | Train: 64 | Validation: 64 | Original demos: 5

Actual training loss (logged optimization batches)



Clean validation loss



Training: train_losses only, exactly as logged; no smoothing and no train_eval_losses substitution. Curves average available seeds at each recorded step.