


countdown_dataset_size_v1 / clean /n1024_f0.0 | countdown_mc


allenai/Olmo-3-1025-7B | wrong target no explicit context | Corruption: incorrect choice | Mode: deterministic | Scope: target | Train: 1024 | Validation: 256 | Original

demos: 5 | Corrupted fraction: 0.0 | Validation context: clean demonstrations

Actual training loss (logged optimization batches)



Clean validation loss



Training: train_losses only, exactly as logged; no smoothing and no train_eval_losses substitution. Curves average available seeds at each recorded step.