



# word\_dataset\_size\_v1 / random / n1024\_f1.0 | compositional:copy\_lower

allenai/Olmo-3-1025-7B | wrong target no explicit context | Corruption: rotation | Mode: random | Scope: target | Train: 1024 | Validation: 256 | Original demos: 5 | Corrupted  
fraction: 1.0 | Validation context: clean demonstrations

Actual training loss (logged optimization batches)



Clean validation loss



Training: train\_losses only, exactly as logged; no smoothing and no train\_eval\_losses substitution. Curves average available seeds at each recorded step.