



word_dataset_size_v1 / random / n64_f1.0 | compositional:gerund_last

allenai/Olmo-3-1025-7B | wrong target no explicit context | Corruption: rotation | Mode: random | Scope: target | Train: 64 | Validation: 256 | Original demos: 5 | Corrupted
fraction: 1.0 | Validation context: clean demonstrations

Actual training loss (logged optimization batches)



Clean validation loss



Checkpoint
stage1-step0 stage1-step1000 stage1-step2000

Training: train_losses only, exactly as logged; no smoothing and no train_eval_losses substitution. Curves average available seeds at each recorded step.